

EDA using R

```
# Calling the Libraries used
library(readxl)
library(dplyr)
library(lubridate)
library(readr)
library(ggplot2)
library(ggthemes)
library(tidyr)
library(DT)
library(scales)
library(stringr)
library(knitr)
library(FactoMineR)
library(ggpubr)
library(kableExtra)
library(magrittr)
library(ggfortify)
```

```
# Reading the data from CSV file and loading into a Dataframe
ozone_df <- read.csv("daily_ozone_2022.csv")
# Reading few datasamples from the CSV file
ozone_df <- sample_n(ozone_df, 5000)
# Printing the Dimensions of the Dataframe
dim(ozone_df)
```

```
## [1] 5000 29
```

```
# Examining the structure of the Dataframe
str(ozone_df)
```

```
## 'data.frame': 5000 obs. of 29 variables:
## $ State.Code : int 51 6 21 32 53 48 51 6 32 22 ...
## $ County.Code : int 113 77 185 33 73 355 113 25 3 57 ...
## $ Site.Num : int 3 1003 4 101 5 25 3 1003 299 4 ...
## $ Parameter.Code : int 44201 44201 44201 44201 44201 44201 44201 44201 44201 44201 ...
## $ POC : int 1 1 1 1 1 2 1 1 1 1 ...
## $ Latitude : num 38.5 38 38.4 39 49 ...
## $ Longitude : num -78.4 -121.3 -85.4 -114.2 -122.6 ...
## $ Datum : chr "WGS84" "WGS84" "WGS84" "WGS84" ...
## $ Parameter.Name : chr "Ozone" "Ozone" "Ozone" "Ozone" ...
## $ Sample.Duration : chr "8-HR RUN AVG BEGIN HOUR" "8-HR RUN AVG BEGIN HOUR" "8-HR RUN AVG BEGIN HOUR" ...
## $ Pollutant.Standard : chr "Ozone 8-hour 2015" "Ozone 8-hour 2015" "Ozone 8-hour 2015" "Ozone 8-hour 2015" ...
## $ Date.Local : chr "2022-10-29" "2022-05-12" "2022-10-18" "2022-10-29" ...
## $ Units.of.Measure : chr "Parts per million" "Parts per million" "Parts per million" "Parts per million" ...
```

```
## $ Event.Type      : chr "None" "None" "None" "None" ...
## $ Observation.Count : int 17 3 17 17 17 17 17 17 17 17 ...
## $ Observation.Percent : num 100 18 100 100 100 100 100 100 100 100 ...
## $ Arithmetic.Mean   : num 0.0385 0.0387 0.0218 0.0398 0.0233 ...
## $ X1st.Max.Value    : num 0.042 0.04 0.023 0.042 0.04 0.011 0.044 0.034 0.028 0.035 ...
## $ X1st.Max.Hour     : int 20 9 12 20 10 21 18 9 9 9 ...
## $ AQI               : int 39 37 21 39 37 10 41 31 26 32 ...
## $ Method.Code       : int NA NA NA NA 87 87 47 NA 87 NA ...
## $ Method.Name       : chr " - " " - " " - " " - " ...
## $ Local.Site.Name   : chr "Shenandoah NP - Big Meadows" "Stockton - University Park" "BUCKNER" "G
## $ Address           : chr "SHENANDOAH NP BIG MEADOWS" "702 N Aurora Street, Stockton, CA 95202" "
## $ State.Name        : chr "Virginia" "California" "Kentucky" "Nevada" ...
## $ County.Name       : chr "Madison" "San Joaquin" "Oldham" "White Pine" ...
## $ City.Name         : chr "Not in a city" "Stockton" "Buckner" "Not in a city" ...
## $ CBSA.Name         : chr "" "Stockton-Lodi, CA" "Louisville/Jefferson County, KY-IN" "" ...
## $ Date.of.Last.Change : chr "2023-03-16" "2023-03-17" "2023-02-20" "2023-03-16" ...
```

```
# Examining the Dataframe, it appears some columns should be converted to Factors and dates
ozone_df$Parameter.Code <- as.factor(ozone_df$Parameter.Code)
ozone_df$Units.of.Measure <- as.factor(ozone_df$Units.of.Measure)
# Handling Date columns
ozone_df$Date.Local <- as.Date(ozone_df$Date.Local)
ozone_df$Date.of.Last.Change <- as.Date(ozone_df$Date.of.Last.Change)
# Printing sample rows from Dataframe
kbl(head(ozone_df[1:6,c(1:8)]), caption = "Ozone Pollutant Data",booktabs = T) %>% kable_styling(latex_
```

Table 1: Ozone Pollutant Data


State.Code	County.Code	Site.Num	Parameter.Code	POC	Latitude	Longitude	Datum
51	113	3	44201	1	38.52310	-78.43471	WGS84
6	77	1003	44201	1	37.96158	-121.28141	WGS84
21	185	4	44201	1	38.40020	-85.44428	WGS84
32	33	101	44201	1	39.00512	-114.21593	WGS84
53	73	5	44201	1	48.95074	-122.55441	WGS84
48	355	25	44201	2	27.76534	-97.43426	WGS84

```
# Checking for Nulls in the Ozone Dataframe
colSums(is.na(ozone_df))
```

```
## State.Code      County.Code      Site.Num      Parameter.Code
## 0                0                0                0
## POC              Latitude          Longitude        Datum
## 0                0                0                0
## Parameter.Name   Sample.Duration  Pollutant.Standard Date.Local
## 0                0                0                0
## Units.of.Measure Event.Type       Observation.Count Observation.Percent
## 0                0                0                0
## Arithmetic.Mean  X1st.Max.Value  X1st.Max.Hour    AQI
## 0                0                0                0
## Method.Code      Method.Name      Local.Site.Name   Address
```


```
##           1743           0           0           0
##      State.Name      County.Name      City.Name      CBSA.Name
##           0           0           0           0
## Date.of.Last.Change
##           0
```

```
# Plotting Histogram to examine the distribution of AQI
ggplot(ozone_df,aes(x=AQI))+geom_histogram(aes(color="red"))+ggtitle(label = 'Histogram of AQI')+ theme
```



Histogram Results: The Histogram results indicate that the AQI is heavily distributed between 20 and 65 and has a long tail on the right indicating outliers.

```
# Plotting Boxplot to examine the distribution of AQI
ggplot(ozone_df,aes(y=AQI))+geom_boxplot(aes(color="red"))+ggtitle(label = 'Boxplot of AQI')+ theme(leg
```



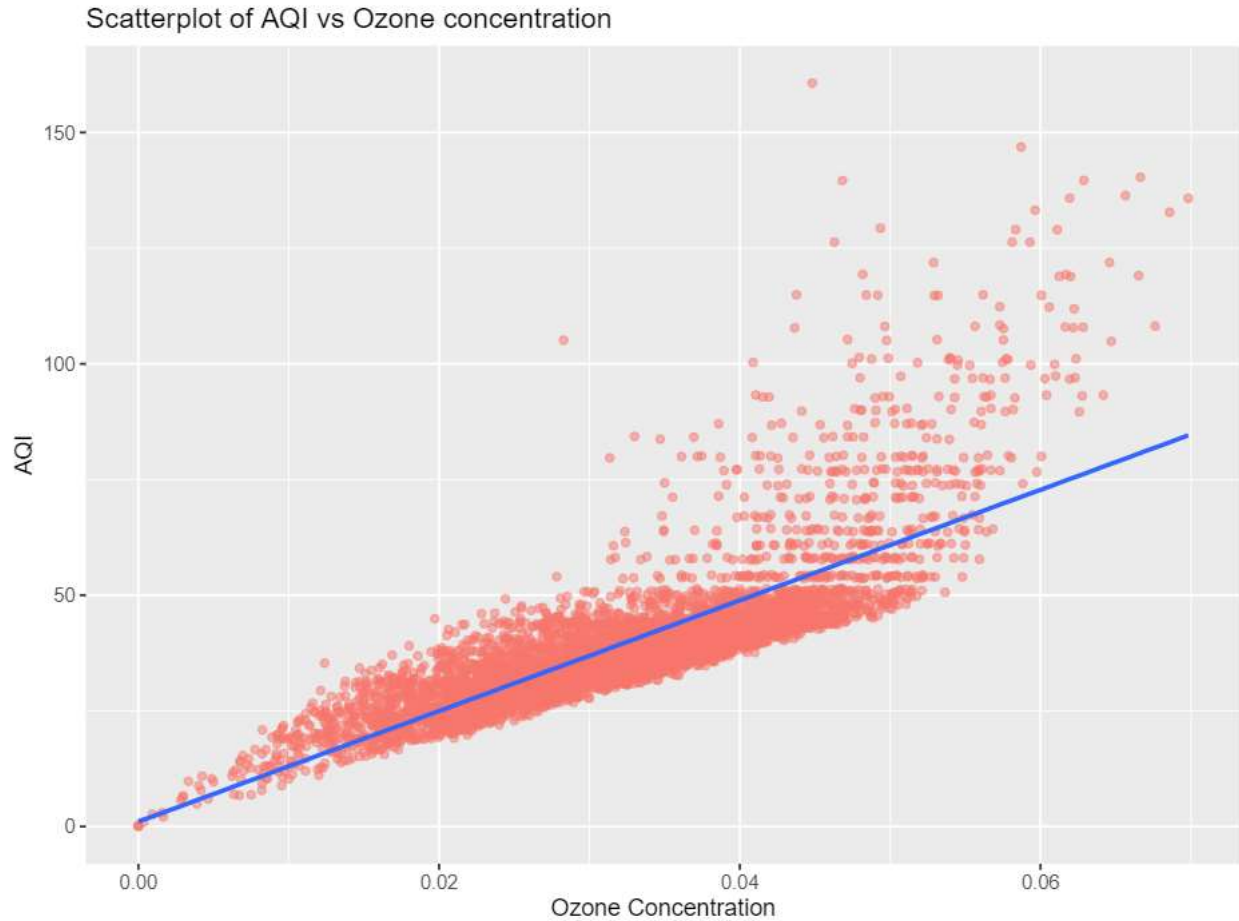
Boxplot Results: The Boxplot results indicate that the presence of Outliers above an AQI of 65 and also below 20. Most of the Data is between 20 and 65.

```
# Computing the first and third quartiles
q1 <- quantile(ozone_df$AQI, 0.25)
q3 <- quantile(ozone_df$AQI, 0.75)
iqr <- q3 - q1
# Calculate the lower and upper cutoffs for outliers
lower <- q1 - 1.5 * iqr
upper <- q3 + 1.5 * iqr
# Filter AQI to find outliers
AQI_outliers <- ozone_df %>%
  filter(AQI > upper | AQI < lower)
# Printing the top few rows of Outliers
kbl(head(AQI_outliers[1:6,c(1:8)]), caption = "Outliers in the AQI",booktabs = T) %>% kable_styling(lat

ggplot(ozone_df,aes(x=Arithmetic.Mean,y=AQI))+geom_point(aes(color="Blue"),alpha=0.5,position="jitter")
```


Table 2: Outliers in the AQI

State.Code	County.Code	Site.Num	Parameter.Code	POC	Latitude	Longitude	Datum
48	355	25	44201	2	27.76534	-97.43426	WGS84
6	43	6	44201	1	37.54377	-119.83957	NAD83
48	201	1034	44201	2	29.76800	-95.22058	WGS84
15	3	1004	44201	2	21.30338	-157.87117	WGS84
39	167	4	44201	1	39.43212	-81.46044	NAD83
4	21	3003	44201	1	32.95436	-111.76225	WGS84



Scatterplot Results: The Scatterplot indicates a strong positive relation between the Ozone Concentration and AQI.

```
# Calculating the correlation coefficient
cor(ozone_df$Arithmetic.Mean,ozone_df$AQI)
```

```
## [1] 0.8041176
```

Correlation Results: The output of the correlation coefficient indicates a Strong positive correlation (0.8) between the Ozone concentration and the Air Quality Index values.

Data Preparation and Building Models using R

```
# Calling the Libraries used
```

```
library(readxl)
library(dplyr)
library(lubridate)
library(readr)
library(ggplot2)
library(ggthemes)
library(tidyr)
library(scales)
library(stringr)
library(knitr)
library(FactoMineR)
library(ggpubr)
library(kableExtra)
library(magrittr)
library(ggfortify)
library(visdat)
library(usmap)
library(plotly)
library(leaflet)
library(magrittr)
library(treemap)
library(olsrr)
```

```
# Function to create a dataframe by reading the contents from csv file
```

```
create_dataframe <- function(df,file_name){
  df <- read.csv(file_name,stringsAsFactors=FALSE)
  return(df)
}
```

```
# Function to rename the columns into required format
```

```
rename_columns <- function(df){
  df <- df %>% select(state_code=State.Code,county_code=County.Code,site_num=Site.Num,
                     date_local=Date.Local,
                     new_col=Arithmetic.Mean,AQI
                     )
  return (df)
}
```

```
# Function to format the columns by changing the datatype to factors and Dates
```

```
format_columns <- function(df){
  df$state_code <- as.factor(df$state_code)
  df$county_code <- as.factor(df$county_code)
  df$site_num <- as.factor(df$site_num)
  df$date_local <- as.Date(df$date_local)
  df$AQI[is.na(df$AQI)] <- 0
  return (df)
}
```

```

# Function to find nulls in the dataframe by taking the column sums
find_nulls <- function(df){
  colSums(is.na(df))
}

# Function to find the correlation between the AQI and the field new_col
find_AQI_corr <- function(df){
  cor(df$new_col,df$AQI)
}

# Function to remove duplicates by grouping the data on common columns and retrieving only the distinct
remove_duplicates <- function(df) {
  df <- df %>% group_by(state_code,county_code,site_num,date_local) %>% distinct(state_code,county_code,site_num,date_local)
  return (df)
}

```

```

# Reading the contents of NO2 concentration into a dataframe for performing the EDA analysis
NO2_EDA_df <- create_dataframe(NO2_EDA_df,"daily_NO2_2022.csv")
# Selecting the required columns from the Dataframe
NO2_EDA_df <- NO2_EDA_df %>% select(state_code=State.Code,county_code=County.Code,site_num=Site.Num,
                                   date_local=Date.Local,
                                   NO2=Arithmetic.Mean,AQI,
                                   state_name=State.Name,county_name=County.Name,city_name=City.Name)
# Converting the columns datatype to Factor and Date
NO2_EDA_df <- format_columns(NO2_EDA_df)

```

```

# This function calls multiple functions to prepare the data for merging
data_preparation <- function(df,filename){
  # Calling function to create dataframe
  df <- create_dataframe(df,filename)
  # Calling function to rename the columns
  df <- rename_columns(df)
  # Function to Format the columns
  df <- format_columns(df)
  # Function to remove duplicates
  df <- remove_duplicates(df)
  return(df)
}

```

```

# Calling Data preparation function on multiple source files each resulting in individual dataframes
NO2_df <- data_preparation(df=NO2_df,filename="daily_NO2_2022.csv")
ozone_df <- data_preparation(ozone_df,"daily_ozone_2022.csv")
SO2_df <- data_preparation(SO2_df,"daily_SO2_2022.csv")
CO_df <- data_preparation(CO_df,"daily_CO_2022.csv")
df_81102 <- data_preparation(df_81102,"daily_81102_2022.csv")
wind_df <- data_preparation(wind_df,"daily_WIND_2022.csv")
temp_df <- data_preparation(CO_df,"daily_TEMP_2022.csv")
press_df <- data_preparation(CO_df,"daily_PRESS_2022.csv")
rh_dp_df <- data_preparation(CO_df,"daily_RH_DP_2022.csv")

```

```

# This function renames the "new_col" column into the desired name
rename_col2 <- function(df,col_name){
  df[[col_name]] <- df$new_col
  df <- df %>% select(-new_col)
  return(df)
}

```



```

}
# Function to join dataframes using the "left join"
merge_dataframes <- function(df1,df2,col_name){
  # Performs left join and selects required columns
  merge_df <- left_join(df1,df2,by=c("state_code","county_code","site_num","date_local")) %>% rename(
  # Replaces the nulls in numeric columns after left join with median values
  merge_df[[col_name]][is.na(merge_df[[col_name]])]<-median(merge_df[[col_name]],na.rm=TRUE)
  return(merge_df)
}

```

```

# This section calls the rename_col2 function to rename the "new_col" field with the desired name
ozone_df = rename_col2(ozone_df,col_name="ozone")
SO2_df = rename_col2(SO2_df,col_name="SO2")
NO2_df = rename_col2(NO2_df,col_name="NO2")
CO_df = rename_col2(CO_df,col_name="CO")
df_81102 = rename_col2(df_81102,col_name="PM")
wind_df = rename_col2(wind_df,col_name="wind")
temp_df = rename_col2(temp_df,col_name="temp")
press_df = rename_col2(press_df,col_name="press")
rh_dp_df = rename_col2(rh_dp_df,col_name="RH")

```

```

# Calling the Join function to perform the joins
merge1_df=merge_dataframes(NO2_df,ozone_df,col_name="ozone")
merge2_df=merge_dataframes(merge1_df,SO2_df,col_name="SO2")
merge3_df=merge_dataframes(merge2_df,CO_df,col_name="CO")
merge4_df=merge_dataframes(merge3_df,df_81102,col_name="PM")
merge5_df=merge_dataframes(merge4_df,wind_df,col_name="wind")
merge6_df=merge_dataframes(merge5_df,temp_df,col_name="temp")
merge7_df=merge_dataframes(merge6_df,press_df,col_name="press")
merge8_df=merge_dataframes(merge7_df,rh_dp_df,col_name="RH")

```

```

# Validating nulls in the final merged dataframe
find_nulls(merge8_df)

```

```

## state_code county_code site_num date_local AQI NO2
##      0      0      0      0      0      0
## ozone      SO2      CO      PM      wind      temp
##      0      0      0      0      0      0
## press      RH
##      0      0

```

Visualizing the data

```

# Performing join with the EDA df to retrieve the state name, county name and location details
merge9_df <- merge8_df %>% inner_join(NO2_EDA_df,by=c("state_code","county_code","site_num","date_local"))

```

```

# Including Region details to the dataframe based on the state names
merge10_df <- merge9_df %>% mutate(
  region_name= case_when(
    state_name %in% c("Maine","Vermont","Massachusetts","Rhode Island","Connecticut","New Hampshire")

```



```

state_name %in% c("New York", "Pennsylvania", "New Jersey", "Delaware", "Maryland") ~ "Mid_Atlantic
state_name %in% c("Arkansas", "Louisiana", "Mississippi", "Alabama", "Georgia", "Florida", "Tennessee
state_name %in% c("North Dakota", "South Dakota", "Nebraska", "Kansas", "Missouri", "Iowa", "Minnesot
state_name %in% c("Nevada", "Utah", "Colorado", "Wyoming", "Idaho", "Montana") ~ "Rocky_Mountain_Sta
state_name %in% c("Washington", "California", "Oregon", "Alaska", "Hawaii") ~ "Pacific_coastal",
state_name %in% c("Arizona", "New Mexico", "Oklahoma", "Texas") ~ "South_West",
state_name %in% c("District Of Columbia") ~ "District Of Columbia"
)
)

```

```

# Counting the number of entries in the final dataframe by state.
count_by_state_df <- merge10_df %>% group_by(state_name) %>% count() %>% arrange(desc(n)) %>% rename(ob
count_by_state_df <- head(count_by_state_df, 10)


```

```

# Creating a Bar plot of the count based on the state count
ggplot(data = count_by_state_df, aes(state_name, observations, fill = observations)) + geom_bar(stat = "
geom_text(aes(label = observations), vjust = -1,
size = 3) + scale_fill_gradient(low = "Red", high = "blue") +
labs(x = "", y = "Number of Observations",
title = "US States with the most observations in the dataset") +
scale_y_continuous(labels = comma) + ylim(0, 45000) + theme_bw() + theme(plot.title = element_text(size
angle = 90), axis.text.y = element_text(size = 8),
axis.title = element_text(size = 10),
legend.position = "none")

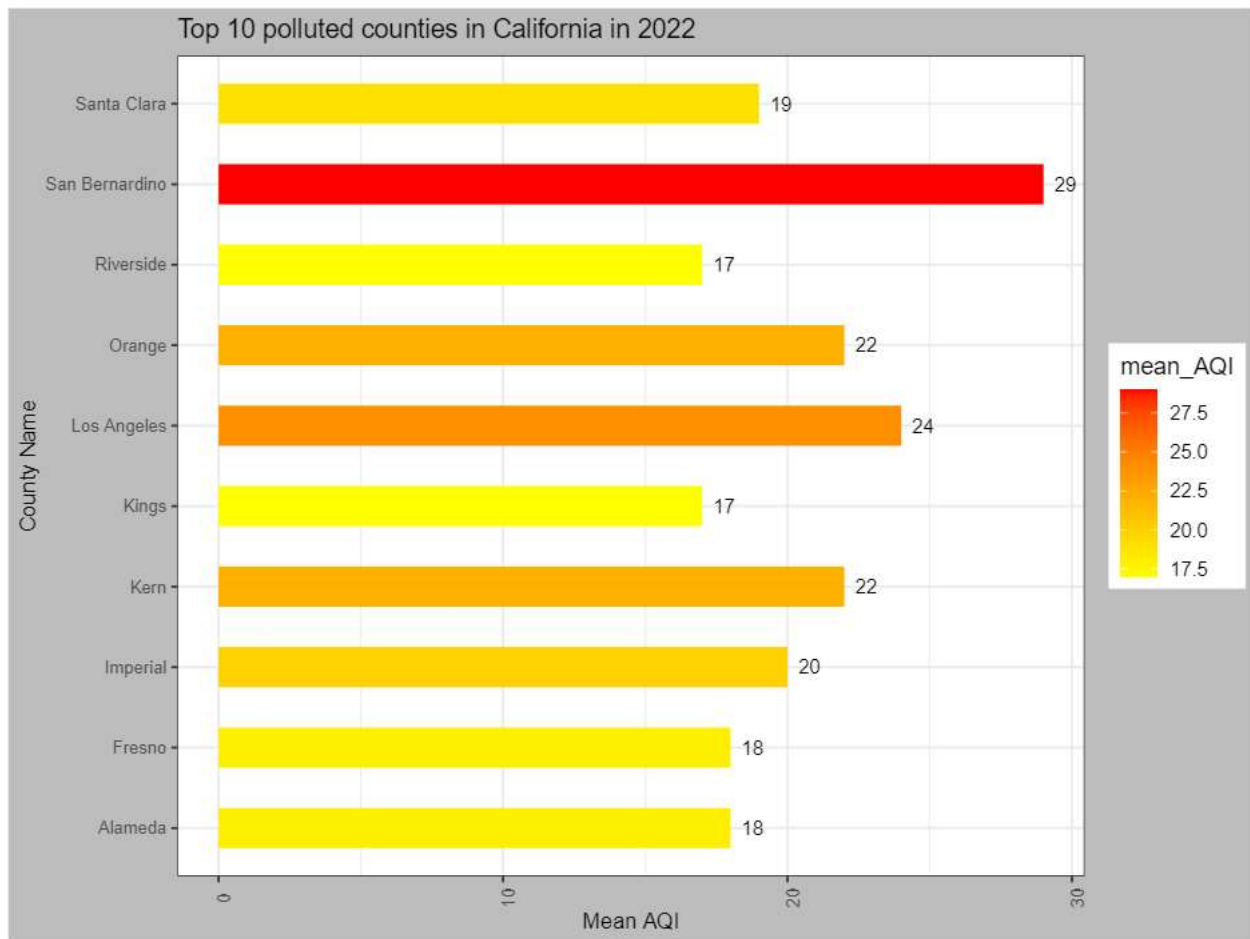
```

US States with the most observations in the dataset



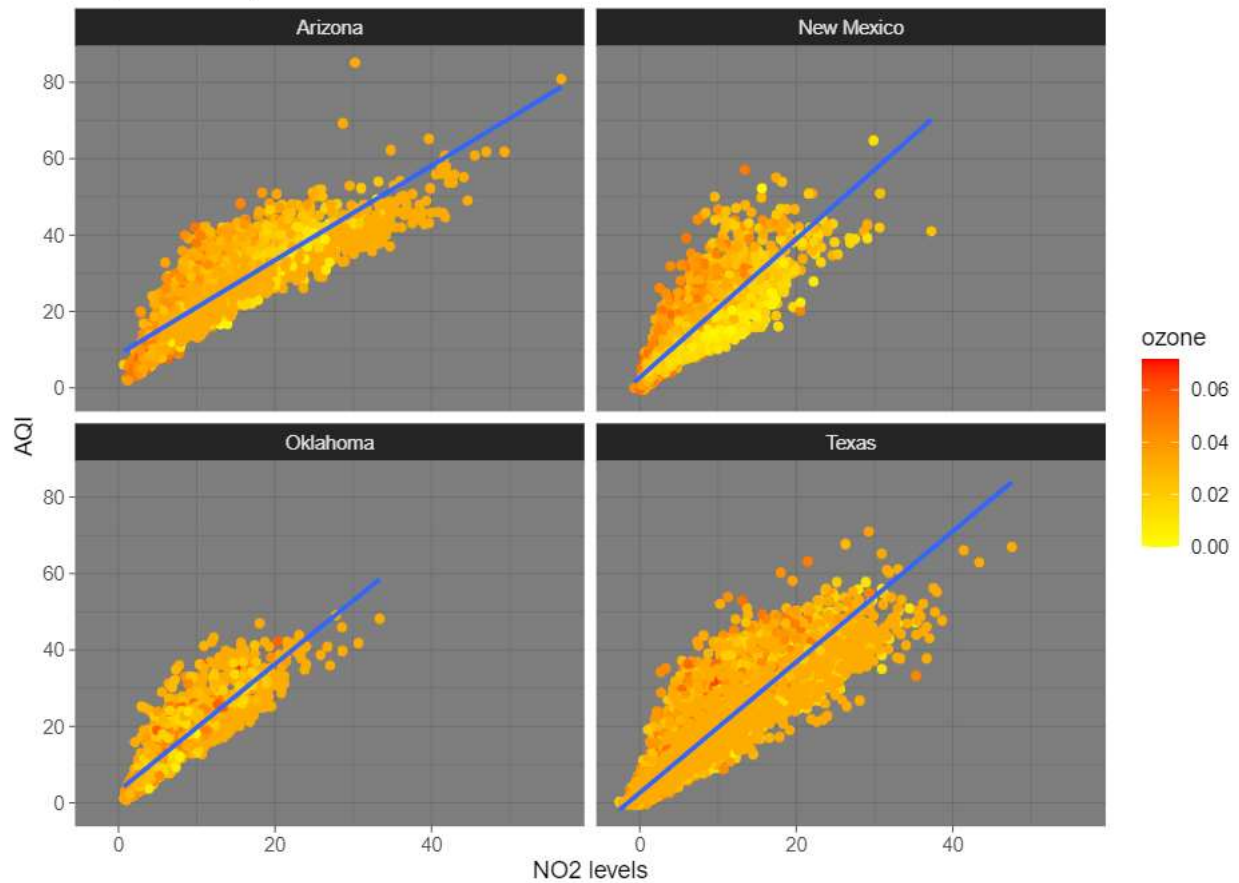
```
# Creating a dataframe for California by identifying the most polluted counties based on the AQI
CA_AQI_df <- merge10_df %>% filter(state_name=="California") %>% group_by(county_name) %>% summarize(me
```

```
# Creating a bar plot of CA_AQI_df with the coordinates flipped for better display
ggplot(data = CA_AQI_df, aes(county_name, mean_AQI)) + geom_bar(stat = "identity",
width = 0.5, aes(fill = mean_AQI)) + scale_fill_gradient(low = "Yellow",
high = "Red") + coord_flip() + labs(x = "County Name",
y = "Mean AQI", title = "Top 10 polluted counties in California in 2022") +
geom_text(aes(label = mean_AQI), hjust = -0.5,
size = 3) + theme_bw() + theme(plot.title = element_text(size = 12),
axis.text.x = element_text(size = 8, angle = 90), axis.text.y = element_text(size = 8),
axis.title = element_text(size = 10), plot.background = element_rect(fill = "Grey"))
```



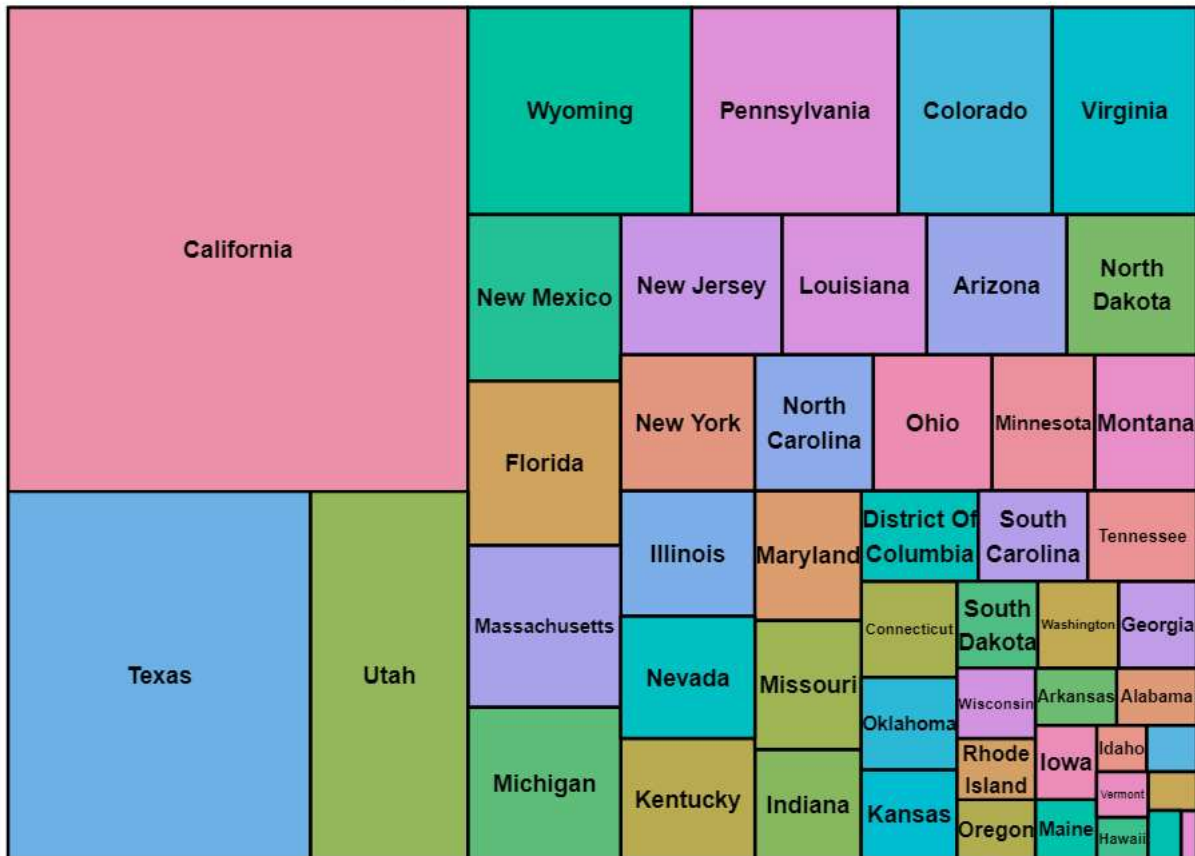
```
# Creating a dataframe for South west region data
SW_df1 <- merge10_df %>% filter(region_name=="South_West")
# Plotting a scatter plot with facet wrap to display the AQI vs NO2 levels based on Ozone concentration
ggplot(SW_df1, aes(x=NO2,y=AQI,color=ozone)) + facet_wrap(~state_name)+
  geom_point(alpha=0.5)+ geom_jitter()+ geom_smooth(method="lm",se=FALSE)+
  scale_color_gradient(low = 'yellow', high = 'red')+
  labs(y="AQI",
       x="NO2 levels",
       title="AQI vs NO2 by ozone levels in South West in 2022")+
  theme_dark()
```


AQI vs NO2 by ozone levels in South West in 2022



```
# Creating a treemap of number of entries in the Ozone dataset
count_df <- merge10_df %>% group_by(state_name) %>% count()
treemap( count_df, index = "state_name", vSize = "n", type = "index", title = "Treemap of Ozone Observa
```

Treemap of Ozone Observations in each state



```
# Calculating the Median value of NO2 in the dataset across all states
median_NO2 <- median(merge10_df$NO2)
median_NO2
```

```
## [1] 5.85
```

```
# Creating a compare dataframe that compares the National median versus each state value
compare_df1 <- merge10_df %>% group_by(state_name) %>% summarize(median_state_NO2=median(NO2)) %>% mutate(difference=median_state_NO2-NO2)
compare_df1 <- head(compare_df1,20)
```

```
# Plotting the compare dataframe displaying each state with above and below the national median values
ggplot(compare_df1, aes(x=state_name, y=difference)) +
  geom_bar(stat='identity', aes(fill=compare_val), width=.5) +
  scale_fill_manual(name="NO2 Median Values",
                    labels = c("Above National Average", "Below National Average"),
                    values = c("above"="#00ba38", "below"="#f8766d")) +
  coord_flip()+theme_bw()+
  theme(plot.title = element_text(size=12),axis.text.x= element_text(size=8),axis.text.y= element_text(size=8),
        y="Difference from National Average",
        title="NO2 Median Concentration - National vs State Average")
```



Model Building

```
# Standardizing the values to build the model using the scale function
merge8_df$NO2 <- scale(merge8_df$NO2)
merge8_df$Ozone <- scale(merge8_df$Ozone)
merge8_df$SO2 <- scale(merge8_df$SO2)
merge8_df$CO <- scale(merge8_df$CO)
merge8_df$wind <- scale(merge8_df$wind)
merge8_df$temp <- scale(merge8_df$temp)
merge8_df$press <- scale(merge8_df$press)
merge8_df$RH <- scale(merge8_df$RH)
```

```
# Creating training and test datasets by splitting them in 75:25 ratio
gp <- runif(nrow(merge8_df))
train_df <- merge8_df[gp < 0.75, ]
test_df <- merge8_df[gp >= 0.75, ]
# Printing the rows in the train and test sets
nrow(train_df)
```

```
## [1] 116464
```

```
nrow(test_df)
```

```
## [1] 39004
```



```
# Creating a linear regression model on the train dataset
model1 <- lm(AQI ~ NO2+ozone+SO2+PM+CO+wind+temp+press+RH+date_local
, data=train_df )
summary(model1)
```

```
##
## Call:
## lm(formula = AQI ~ NO2 + ozone + SO2 + PM + CO + wind + temp +
##     press + RH + date_local, data = train_df)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -27.727  -2.960  -1.008   2.157   78.442
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  92.6116876   2.6086686   35.502 < 2e-16 ***
## NO2          10.7033274   0.0161371  663.273 < 2e-16 ***
## ozone         0.1098810   0.0154328    7.120 1.09e-12 ***
## SO2          -0.0909903   0.0144301   -6.306 2.88e-10 ***
## PM           0.0275975   0.0013108   21.054 < 2e-16 ***
## CO           -0.5877026   0.0151433  -38.810 < 2e-16 ***
## wind         -0.1810164   0.0139879  -12.941 < 2e-16 ***
## temp          0.0410445   0.0147863    2.776 0.00551 **
## press        -0.0870037   0.0140760   -6.181 6.39e-10 ***
## RH           -0.6743306   0.0146805  -45.934 < 2e-16 ***
## date_local   -0.0040564   0.0001361  -29.814 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 4.766 on 116453 degrees of freedom
## Multiple R-squared:  0.8305, Adjusted R-squared:  0.8304
## F-statistic: 5.704e+04 on 10 and 116453 DF,  p-value: < 2.2e-16
```

```
# Using the backward fit method to identify the features that can be excluded from the model
backwardfit.p<-ols_step_backward_p(model1,prem=.05)
backwardfit.p
```

```
## [1] "No variables have been removed from the model."
```

```
# Finding the best model using the ols_step_best_subset function
modcompare<-ols_step_best_subset(model1)
modcompare
```

```
##                               Best Subsets Regression
## -----
## Model Index    Predictors
## -----
##      1         NO2
##      2         NO2 RH
##      3         NO2 CO RH
##      4         NO2 CO RH date_local
```

```
##      5      NO2 PM CO RH date_local
##      6      NO2 PM CO wind RH date_local
##      7      NO2 ozone PM CO wind RH date_local
##      8      NO2 ozone SO2 PM CO wind RH date_local
##      9      NO2 ozone SO2 PM CO wind press RH date_local
##     10      NO2 ozone SO2 PM CO wind temp press RH date_local
```

```
## -----
##
##                                     Subsets Regression Summary
## -----
```

## Model	R-Square	Adj. R-Square	Pred R-Square	C(p)	AIC	SBIC	SBC
## 1	0.8210	0.8210	0.821	6456.8733	700507.8046	369997.0744	700536.8006
## 2	0.8256	0.8256	0.8256	3328.8701	697505.7067	366995.0240	697544.3680
## 3	0.8279	0.8279	0.8277	1772.5573	695982.6800	365472.0455	696031.0067
## 4	0.8292	0.8292	0.829	847.7867	695068.1444	364557.5582	695126.1364
## 5	0.8300	0.8300	0.8298	319.1738	694542.1261	364031.5798	694609.7834
## 6	0.8302	0.8302	0.83	150.1134	694373.3907	363862.8597	694450.7134
## 7	0.8303	0.8303	0.8301	93.0517	694316.3830	363805.8577	694403.3710
## 8	0.8304	0.8304	0.8302	53.4082	694276.7595	363766.2394	694373.4129
## 9	0.8304	0.8304	0.8302	16.7054	694240.0620	363729.5479	694346.3807
## 10	0.8305	0.8304	0.8303	11.0000	694234.3562	363723.8436	694350.3402

```
## -----
## AIC: Akaike Information Criteria
## SBIC: Sawa's Bayesian Information Criteria
## SBC: Schwarz Bayesian Criteria
## MSEP: Estimated error of prediction, assuming multivariate normality
## FPE: Final Prediction Error
## HSP: Hocking's Sp
## APC: Amemiya Prediction Criteria
```

```
# Building a final model based on the above findings and printing the summary
model2 <- lm(AQI ~ NO2+ozone+SO2+PM+CO+wind+temp+press+RH+date_local
, data=train_df )
summary(model2)
```

```
##
## Call:
## lm(formula = AQI ~ NO2 + ozone + SO2 + PM + CO + wind + temp +
##      press + RH + date_local, data = train_df)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -27.727  -2.960  -1.008   2.157  78.442
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  92.6116876   2.6086686   35.502  < 2e-16 ***
## NO2          10.7033274   0.0161371  663.273  < 2e-16 ***
## ozone         0.1098810   0.0154328    7.120 1.09e-12 ***
## SO2          -0.0909903   0.0144301   -6.306 2.88e-10 ***
## PM           0.0275975   0.0013108   21.054  < 2e-16 ***
## CO          -0.5877026   0.0151433  -38.810  < 2e-16 ***
```

```
## wind      -0.1810164  0.0139879 -12.941 < 2e-16 ***
## temp      0.0410445  0.0147863   2.776  0.00551 **
## press     -0.0870037  0.0140760  -6.181  6.39e-10 ***
## RH        -0.6743306  0.0146805 -45.934 < 2e-16 ***
## date_local -0.0040564  0.0001361 -29.814 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 4.766 on 116453 degrees of freedom
## Multiple R-squared:  0.8305, Adjusted R-squared:  0.8304
## F-statistic: 5.704e+04 on 10 and 116453 DF,  p-value: < 2.2e-16
```

```
# Adding a new column for prediction in the Training and test dataset
train_df$pred_AQI <- predict(model2 ,train_df)
test_df$pred_AQI <- predict(model2 ,test_df)
```

```
# Function to calculate the R-squared values manually
r_squared <- function(predcol, ycol) {
  tss = sum( (ycol - mean(ycol))^2 )
  rss = sum( (predcol - ycol)^2 )
  1 - rss/tss
}
# Function to calculate the Root mean square error values manually
rmse <- function(predcol, ycol) {
  res = predcol-ycol
  sqrt(mean(res^2))
}
# Calculating the RMSE for training and test datasets
rmse_train <- rmse(train_df$pred_AQI,train_df$AQI)
sprintf("The RMSE value of Training Dataset is %s", round(rmse_train,2) )
```

```
## [1] "The RMSE value of Training Dataset is 4.77"
```

```
rmse_test <- rmse(test_df$pred_AQI,test_df$AQI)
sprintf("The RMSE value of Training Dataset is %s", round(rmse_test,2) )
```

```
## [1] "The RMSE value of Training Dataset is 4.77"
```

```
# Evaluate the r-squared on both training and test data.and print them
rsq_train <- r_squared(train_df$pred_AQI,train_df$AQI)
sprintf("The R-squared value of Training Dataset is %s", round(rsq_train,2) )
```

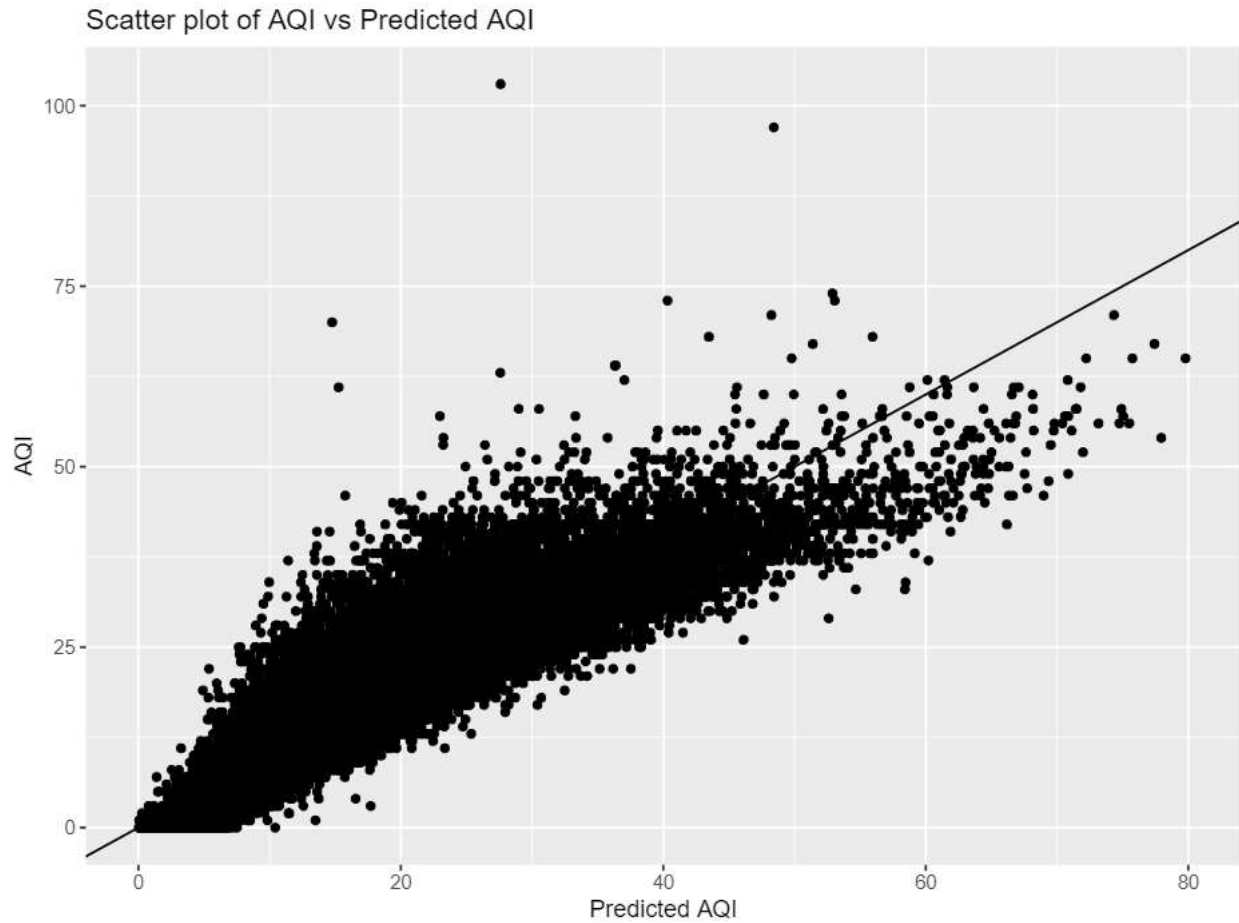
```
## [1] "The R-squared value of Training Dataset is 0.83"
```

```
rsq_test <- r_squared(test_df$pred_AQI,test_df$AQI)
sprintf("The R-squared value of Test Dataset is %s", round(rsq_test,2) )
```

```
## [1] "The R-squared value of Test Dataset is 0.83"
```



```
# Plot the predictions (on the x-axis) against the outcome (AQI) on the test data
ggplot(test_df, aes(x = pred_AQI, y = AQI)) +
  geom_point() + xlim(0,80)+
  geom_abline()+labs(title="Scatter plot of AQI vs Predicted AQI",x="Predicted AQI", y="AQI")
```



Milestone4_Using_Python

February 11, 2024

```
[1]: # Importing required libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
from sklearn.metrics import mean_squared_error
from math import sqrt

# Ignore warnings
import warnings
warnings.filterwarnings('ignore')

#For building ML model
from sklearn.model_selection import train_test_split

#Different Regressors for ML model
from sklearn import linear_model
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor, \
    AdaBoostRegressor
from sklearn.gaussian_process import GaussianProcessRegressor
from sklearn.gaussian_process.kernels import DotProduct, WhiteKernel, RBF
from sklearn.neighbors import KNeighborsRegressor
from sklearn.neural_network import MLPRegressor
from sklearn.svm import SVR
from sklearn.tree import DecisionTreeRegressor
from sklearn.linear_model import LinearRegression

#For model evaluation
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.model_selection import cross_val_predict
from sklearn.metrics import accuracy_score, precision_score, recall_score, \
    f1_score

[2]: # Creating Dataframe from the merged data
df_pollutants = pd.read_csv('merge10_df.csv')
df_pollutants
```

```

[2]:      Unnamed: 0  state_code  county_code  site_num  date_local  AQI  \
0          1          1          73          23  2022-01-01    1
1          2          1          73          23  2022-01-02    4
2          3          1          73          23  2022-01-03    2
3          4          1          73          23  2022-01-04    8
4          5          1          73          23  2022-01-05   27
...
157739    157740      72          25          7  2022-09-14    5
157740    157741      72          25          7  2022-09-15    6
157741    157742      72          25          7  2022-09-16    8
157742    157743      72          25          7  2022-09-17    5
157743    157744      72          25          7  2022-09-18    2

      NO2      ozone      SO2      CO  PM      wind      temp  \
0    1.308333  0.024765 -0.191667  0.100000  14  21.525000  73.083333
1    1.954167  0.017824 -0.237500  0.145833   7  11.359091  55.141667
2    1.530000  0.031882 -0.380000  0.200000  18  49.100000  31.560000
3    5.455556  0.025353  0.555556  0.211111  18  12.266667  47.493333
4   16.493750  0.012563  0.550000  0.406667  21   5.995833  47.887500
...
157739    3.466667  0.031882  0.345833  0.612500  18   3.729167  60.333333
157740    4.012500  0.031882  0.345833  0.600000  18   3.729167  60.333333
157741    4.575000  0.031882  0.345833  0.704167  18   3.729167  60.333333
157742    3.225000  0.031882  0.345833  0.458333  18   3.729167  60.333333
157743    1.580000  0.031882  0.345833  0.400000  18   3.729167  60.333333

      press      RH  state_name  county_name  city_name  \
0    988.595833  56.191666    Alabama    Jefferson    Birmingham
1    991.304167  56.191666    Alabama    Jefferson    Birmingham
2   1002.450000  56.191666    Alabama    Jefferson    Birmingham
3   1001.420000  56.191666    Alabama    Jefferson    Birmingham
4    996.404167  56.191666    Alabama    Jefferson    Birmingham
...
157739    985.579167  56.191666  Puerto Rico    Caguas    Caguas
157740    985.579167  56.191666  Puerto Rico    Caguas    Caguas
157741    985.579167  56.191666  Puerto Rico    Caguas    Caguas
157742    985.579167  56.191666  Puerto Rico    Caguas    Caguas
157743    985.579167  56.191666  Puerto Rico    Caguas    Caguas

      region_name
0    South_East
1    South_East
2    South_East
3    South_East
4    South_East
...
157739    NaN

```



```

157740      NaN
157741      NaN
157742      NaN
157743      NaN

```

```
[157744 rows x 19 columns]
```

```

[3]: # Defining the columns to select
y = df_pollutants['AQI']
columns_to_select = ['NO2', 'ozone', 'SO2', 'CO', 'PM', 'wind', 'temp', 'RH']
x = df_pollutants[columns_to_select]

#Split data into test and training sets
X_train, X_test, y_train, y_test = train_test_split(x, y, test_size = 0.2,
↳random_state = 42)
X_train.shape,X_test.shape,y_train.shape,y_test.shape

```

```
[3]: ((126195, 8), (31549, 8), (126195,), (31549,))
```

```

[4]: def model_assess(X_train, X_test, y_train, y_test, model, title ):
      """
      This function will be used to build the model. It takes train and test_
↳attributes as the
      input and returns model parameters as output
      """
      model.fit(X_train, y_train)
      y_train_pred = model.predict(X_train)
      y_test_pred  = model.predict(X_test)

      train_mse = mean_squared_error(y_train, y_train_pred)
      train_r2 = r2_score(y_train, y_train_pred)
      test_mse = mean_squared_error(y_test, y_test_pred)
      test_r2 = r2_score(y_test, y_test_pred)

      r_squared = r2_score(y_test_pred,y_test)
      accuracy = round(r_squared*100,2)

      result = [str(title),test_mse, test_r2,r_squared,accuracy]

      return result

```

```

[5]: # Creating a list of algorithms to perform the testing
algs = [LinearRegression(),
        KNeighborsRegressor(),
        RandomForestRegressor(),
        DecisionTreeRegressor(max_features = 'auto', max_depth=3,
↳random_state=42),

```

```
GradientBoostingRegressor(n_estimators=100, max_depth=3,
↪random_state=42)]
```

```
[6]: # Creating a list where the results will be captured
results_list=[]
# Looping through each algorithm and build the model
for alg in algs:
    results_list.append(model_assess(X_train, X_test, y_train, y_test, alg,
↪title = alg) )
# Creating a Dataframe of the Model results
regression_results_df=pd.DataFrame(results_list, columns=['Algorithm','Test_
↪MSE','Test R2','R Squared','Accuracy'])
regression_results_df
```


```
[6]:
```

	Algorithm	Test MSE	Test R2 \
0	LinearRegression()	23.157948	0.826961
1	KNeighborsRegressor()	21.482055	0.839484
2	RandomForestRegressor()	17.442603	0.869667
3	DecisionTreeRegressor(max_depth=3, max_feature...	22.332338	0.833130
4	GradientBoostingRegressor(random_state=42)	17.199274	0.871485

	R Squared	Accuracy
0	0.793633	79.36
1	0.815963	81.60
2	0.854467	85.45
3	0.802207	80.22
4	0.851299	85.13

1 Feature Selection Using Pearson Correlation

```
[33]: #Using Pearson Correlation
plt.figure(figsize=(12,10))
corr_cols=['NO2', 'ozone', 'SO2', 'CO', 'PM', 'wind', 'temp', 'RH']
# Creating a correlation matrix
corr_matrix = df_pollutants[corr_cols].corr().round(2)
sns.heatmap(corr_matrix, annot=True, cmap=plt.cm.Reds)
plt.show()
```



```
[30]: #Correlation with output variable
cor_target = abs(cor["AQI"])
#Selecting highly correlated features
relevant_features = cor_target[cor_target>0.3]
relevant_features
```

```
[30]:      NO2  ozone  SO2    CO   PM  wind  temp  RH
NO2    1.00   NaN   NaN  0.44  NaN   NaN   NaN   NaN
ozone   NaN   1.0   NaN   NaN  NaN   NaN   NaN   NaN
SO2     NaN   NaN   1.0   NaN  NaN   NaN   NaN   NaN
CO      0.44   NaN   NaN  1.00  NaN   NaN   NaN   NaN
PM       NaN   NaN   NaN   NaN  1.0   NaN   NaN   NaN
wind     NaN   NaN   NaN   NaN  NaN   1.0   NaN   NaN
temp     NaN   NaN   NaN   NaN  NaN   NaN   1.0   NaN
RH       NaN   NaN   NaN   NaN  NaN   NaN   NaN   1.0
```

```
[10]: # The features NO2,CO are highly correlated with the output variable AQI,Hence,
      ↪we will drop all other features apart from these
```

```
y = df_pollutants['AQI']
columns_to_select = ['NO2', 'CO']
x = df_pollutants[columns_to_select]

#Split data into test and training sets
X_train, X_test, y_train, y_test = train_test_split(x, y, test_size = 0.2,
      ↪random_state = 42)
X_train.shape,X_test.shape,y_train.shape,y_test.shape
```

```
[10]: ((126195, 2), (31549, 2), (126195,), (31549,))
```

```
[11]: # Creating a list where the results will be captured
reg_results_list1=[]
# Looping through each algorithm and build the model
for model in algs:
    reg_results_list1.append(model_assess(X_train, X_test, y_train, y_test,
      ↪model, title = model) )
```

```
[12]: # Creating a Dataframe of the Model results
df_reg=pd.DataFrame(reg_results_list1, columns=['Algorithm','Test MSE','Test_
      ↪R2','R Squared','Accuracy'])
df_reg
```

```
[12]:
```

	Algorithm	Test MSE	Test R2 \
0	LinearRegression()	23.928997	0.821200
1	KNeighborsRegressor()	22.712246	0.830292
2	RandomForestRegressor()	22.510027	0.831803
3	DecisionTreeRegressor(max_depth=3, max_feature...	22.332338	0.833130
4	GradientBoostingRegressor(random_state=42)	18.987270	0.858125

	R Squared	Accuracy
0	0.785549	78.55
1	0.811406	81.14
2	0.814045	81.40
3	0.802207	80.22
4	0.835796	83.58

2 Feature Reduction by PCA

```
[13]: # Selecting list of columns to build the model
columns_to_select = ['NO2', 'ozone', 'SO2', 'CO', 'PM', 'wind', 'temp', 'RH']
x = df_pollutants[columns_to_select]
```



```
[14]: #Split data into test and training sets
X_train, X_test, y_train, y_test = train_test_split(x, y, test_size = 0.2,
↳ random_state = 42)
X_train.shape,X_test.shape,y_train.shape,y_test.shape
```

```
[14]: ((126195, 8), (31549, 8), (126195,), (31549,))
```

```
[15]: from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
# Standardize feature matrix
scaler_pca = StandardScaler()
# Standardize the feature matrix
features = scaler_pca.fit_transform(X_train)
# Create a PCA that will retain 80 of variance
pca = PCA(n_components=0.80, whiten=True)
# Conduct PCA
features_pca = pca.fit_transform(features)
# Show results
print("Original number of features:", features.shape[1])
print("Reduced number of features:", features_pca.shape[1])
```

Original number of features: 8

Reduced number of features: 6

```
[17]: # Standardize the feature matrix
features_test = scaler_pca.transform(X_test)
# Conduct PCA
features_pca_test = pca.transform(features_test)
# Show results
print("Original number of features:", features_test.shape[1])
print("Reduced number of features:", features_pca_test.shape[1])
X_test.shape[1]
```

Original number of features: 8

Reduced number of features: 6

```
[17]: 8
```

```
[18]: # Create linear regression object
reg_pca = linear_model.LinearRegression()
reg_pca.fit(features_pca,y_train)

# Predicting the output

y_pred_pca = reg_pca.predict(features_pca_test)
y_pred_pca.shape
```

```
[18]: (31549,)
```

```
[19]: from sklearn.metrics import mean_absolute_error,mean_squared_error,r2_score

# Mean Squared Error
mse = mean_squared_error(y_pred_pca,y_test)
# Root Mean Squared Error
rmse = np.sqrt(mse)
# Mean Absolute error
mae = mean_absolute_error(y_pred_pca,y_test)
# R2
r_squared = r2_score(y_pred_pca,y_test)
r_squared
print("Mean Squared Error:", mse)
print("Root Mean Squared Error:", rmse)
print("R-squared:", r_squared)
print("Accuracy",round(r_squared*100,2))
```

Mean Squared Error: 67.58882611476984
Root Mean Squared Error: 8.22124237051614
R-squared: 0.2458928916249895
Accuracy 24.59

```
[26]: def model_assess_pca(features_pca, features_pca_test, y_train, y_test, model,
    title = model):
    """
    This function will be used to build the model. It takes train and test
    attributes as the
    input and returns model parameters as output
    """
    model.fit(features_pca, y_train)
    y_test_pred_pca = model.predict(features_pca_test)
    test_mse = mean_squared_error(y_test, y_test_pred_pca)
    test_r2 = r2_score(y_test, y_test_pred_pca)
    r_squared = r2_score(y_test_pred_pca,y_test)
    accuracy = round(r_squared*100,2)
    result = [str(title),test_mse, test_r2,r_squared,accuracy]
    return result
```

```
[27]: # Creating a list where the results will be captured
results_list_pca=[]
# Looping through each algorithm and build the model
for model in algs:
    results_list_pca.append(model_assess_pca(features_pca, features_pca_test,
    y_train, y_test, model, title = model) )
```

```
[28]: # Create DataFrame to capture the model results
columns = ['Algorithm', 'Test MSE', 'Test R2', 'R Squared', 'Accuracy']
df_regression = pd.DataFrame(results_list_pca, columns=columns)
```

```
df_regression
```

[28]:

	Algorithm	Test MSE	Test R2	\
0	LinearRegression()	67.588826	0.494969	
1	KNeighborsRegressor()	27.964259	0.791048	
2	RandomForestRegressor()	26.463282	0.802263	
3	DecisionTreeRegressor(max_depth=3, max_feature...	61.658604	0.539280	
4	GradientBoostingRegressor(random_state=42)	38.324697	0.713634	

	R Squared	Accuracy
0	0.245893	24.59
1	0.750876	75.09
2	0.753679	75.37
3	0.147110	14.71
4	0.568430	56.84